

Project Report:

Enhanced Image Inpainting using Transformer-GAN Hybrid Models

Group: 32

Team mates:

Veera Saiva Sai Lakshmi Kalyan (50595073)

Taran Mamidala (50604177)

Jupalli Sai Abhisekh Reddy (50604117)

Introduction:

The primary objective of this project is to significantly enhance image inpainting technology by effectively implementing and integrating Transformer architecture and Generative Adversarial Networks (GANs). For this project we have decided to select and work with CelebA dataset due to its excellent attributes that complement the objectives of this project. With 202,599 high-quality facial images, CelebA dataset is very diverse and complex and hence sufficient for providing diversity for enhancing robustness in deep learning models trained for image reconstruction tasks. CelebA's rich collection of facial attributes, expressions, and poses is particularly beneficial for learning detailed and intricate patterns necessary for realistic facial image inpainting. The dataset allows challenging cases, including irregular masks and large missing regions, that mimic realworld inpainting situations. Its prevalent use in past work proves it to be a valid benchmark, which has been continually verified by using metrics like PSNR and SSIM. Taking these strengths forward, this project aims at developing a Transformer-GAN hybrid model that not only enhances the image reconstruction quality but also sets a higher standard for image inpainting solutions.

Dataset and Environment Setup:

We established an organized environment in order to work with and preprocess CelebA efficiently. First, we replicated our compressed version of CelebA dataset (img_align_celeba.zip) from Google Drive and unzip it locally such that all files of images became readily accessible. To ease preprocessing and loading images, we utilized large Python libraries like PyTorch and torchvision for operations on the dataset and image transformations, pandas for metadata and annotation handling, and PIL for image-specific tasks like reading and conversion of formats. We then checked the integrity and completeness of the dataset after extraction through image quantity checks and ensured that metadata files were present. This diligent preparation made it compatible with our deep learning pipeline while keeping data quality in check essential to successful model training.

EDA:

We conducted a thorough Exploratory Data Analysis (EDA) to familiarize and prepare and know the CelebA dataset in-depth before training our hybrid model of Transformer-GAN. We first read and concatenated the dataset features and partition information utilizing pandas. We converted the attribute labels from the original (-1, 1) format to a more natural binary format (0, 1). We also divided the dataset into training (162,770 samples), validation (19,867 samples), and test sets (19,962 samples) as per the official splits of the dataset. To achieve deeper insight, we computed descriptive statistics for all attributes and revealed their distribution and prevalence in the data set. For instance, attributes like 'Attractive', 'Young', and 'Smiling' contained various distributions that were crucial in guiding our model's learning schedule. We also verified whether missing values existed in each column of attributes to ascertain that the data set was whole and valid. Then, we've defined significant changes with torchvision like resizing, random horizontal flipping, and normalizing to prepare images efficiently for deep learning. We created a custom 'CelebADataset' class with proper error handling to only select existing and readable images, and data became more trustworthy. Using PyTorch's 'DataLoader', we prepared data batching efficiently for training, validation, and testing. Furthermore augmenting our data preparation, we introduced random masks to simulate realistic inpainting conditions in real-world scenarios by incorporating random square occlusions into images. This was achieved via the 'InpaintingCelebADataset' class, such that our model experiences realistic training conditions similar to those in real-world applications. Finally, we presented masked images, their corresponding masks, and original images alongside each other, verifying our masking strategy and dataset preparation. This thorough exploratory procedure ensured that our dataset was optimally structured and ready for effective TransformerGAN hybrid model training.

Visualizations:




Figure1




Figure2

We performed further exploratory visualizations to better understand the nature and interrelation within the CelebA dataset. Figure 1 depicts the number of positive samples per facial attribute in the training set, which clearly demonstrates the frequency and attribute imbalance between attributes. Surprisingly, attributes such as "No_Beard" and "Young" contained significantly more instances than attributes such as "Bald," "Mustache," and "Gray_Hair," affecting our perception of potential biases with model training.

In Figure 2, we present a correlation heatmap exploring correlations among the 40 features. By means of this visualization, we are able to identify strong correlations, both positive and negative, enabling further examination of dependencies and redundancies in features. For example, facial hair and gender features possessed stronger correlations, capturing patterns necessary for

wellinformed model design decisions. These graphical examinations were crucial in determining how to construct a good and robust Transformer-GAN hybrid model.

Methodology:

We developed a baseline model of Transformer-GAN hybrid model tailored for the CelebA dataset to solve the image inpainting task efficiently. The model architecture contains two basic components: generator and discriminator.

Generator Architecture:

The generator uses an encoder-decoder structure to effectively capture global information and synthesize high-level textures of images. In the encoder, we utilized successive convolutional layers with ReLU activation and batch normalization to iteratively extract and downsample required features. The decoder reverses the process with transpose convolutional layers to upsample and synthesize the entire image, utilizing Tanh activation to normalize the outputs within the range $[-1, 1]$.

Discriminator Architecture:

Our discriminator is designed to distinguish between model-generated images and real images and give realistic output. It comprises a number of convolutional layers with LeakyReLU activations and batch normalization for efficient feature extraction. After the final convolutional layer, a sigmoid function and adaptive average pooling are employed to produce a confidence score indicating the authenticity of an image.

Training Procedure:

Binary cross-entropy (BCE) loss during adversarial training guided the generator in producing images of realistic output, and one other L1 reconstruction loss as additional pixel fidelity for the reconstructed part. A 0.0002 learning rate Adam optimizer was chosen for good performance due to its apt choice, wherein separate optimizers for the discriminator and generator. Training loop included several epochs of iterative optimisation, which were alternately updated for discriminator and generator. We tracked and printed loss statistics per epoch to monitor convergence as well as stable training.

Results Visualization:

We quantified the visual quality of our inpainted results by examining generated images across several training epochs. The early outputs (Epoch 1) exhibited major reconstruction artifacts and softness in the masked regions. However, by Epoch 2, the detail and texture consistency improvements began to emerge. Additional training up to Epoch 3 and Epoch 4 resulted in progressively clearer and more natural-looking inpainting outputs, which indicates the model's growing capability to reconstruct missing image parts correctly.

After training, we stored the trained generator and discriminator models in final form to ensure reproducibility and simplicity of use in the future. At inference, we loaded the trained generator model to create inpainted images on unseen masked inputs. Inference visualization (Figure: "Inpainted Images (Inference)") confirms the generator's capability in face image restoration, providing improved coherence and realism in reconstruction regions. Apart from that, we also performed a quantitative performance evaluation on the test data. We used the common metrics of Mean Absolute Error (L1 Loss), Peak Signal-to-Noise Ratio (PSNR), and Structural Similarity Index (SSIM). We achieved a mean L1 Loss of 0.0377, representing small reconstruction error. Further, our model achieved a mean PSNR of 23.75 dB, indicating quality restoration, and an SSIM value of 0.7674, which represents high structural fidelity of produced and original images.

Training Images:

Inpainted Images at Epoch 1



Inpainted Images at Epoch 2



Inpainted Images at Epoch 3



Inpainted Images at Epoch 4



Inpainted Images at Epoch 5



Inpainted Images (Inference)



Visualization Graphs & analysis:






Figure3




Figure4




Figure5




Figure6




Figure7

We conducted further quantitative analysis to gain more insight into the performance of our hybrid Transformer-GAN model. Figure 3 depicts the distributions of evaluation metrics across our test set, including L1 Loss, PSNR, and SSIM. The plots of distribution reveal that the majority of images possess low L1 loss values at around 0.0377, uniform PSNR values centered predominantly at 23.75 dB, and high structural similarity with the majority of SSIM values being around 0.7674. These histograms verify our model's consistently good reconstruction performance for different samples. In Figure 4, we present training loss curves over five epochs with the progress of the generator and discriminator monitored. The curves are demonstrating stable and convergent behavior with the discriminative loss decreasing progressively and the generator loss stabilizing after oscillation, signaling an optimally balanced adversarial training process. For more insight into metric correlations, we had plotted scatter plots that indicated correlations between different metrics. Figure 5 is a scatter plot of PSNR vs. SSIM, which indicates a positive correlation between the two variables, meaning that images with higher PSNR values have higher structural similarity. Figure 6 presents scatter plots of L1 Loss vs. PSNR and SSIM separately, which clearly demonstrate inverse correlations since lower L1 losses are highly correlated with higher PSNR and SSIM, which verifies the correspondence between pixel-wise accuracy and perceptual quality. Lastly, Figure 7 provides box plots summarizing the distributions of L1 Loss, PSNR, and SSIM, clearly displaying median values, interquartile ranges, and potential outliers. These graphical

summaries provide a concise view of the variability and stability of the metrics achieved on the entire test dataset.

In summary, these detailed analyses reinforce our model's efficacy in generating high-quality image inpainting results, tested both visually and quantitatively.

Visual Inspection of Test Samples:

We performed a qualitative visual inspection on several test set samples to directly assess the inpainting capability of the Transformer-GAN hybrid model. Figure 8 contains four sample examples, each including the masked input image, the inpainted output created by our model, the original unmasked image, and the absolute pixel-wise difference between the reconstructed image and the original image. The visualization shows that our model effectively recovered visually coherent content in the masked areas, preserving coherence in facial structures, textures, and colors. Discrepancies between original and reconstructed images indicate where the model struggled, e.g., with subtle color differences or small artifacts near boundaries. However, these qualitative results validate our model's strong capability to produce realistic and accurate inpainting, supporting the quantitative results achieved previously.



Extension For Final submission:

As per your request, below is a detailed extension of your report, with the same tone and terminology but now applying it to your three GAN-based models: DCGAN-based Encoder-Decoder model, BigGAN-inspired model with attention, and the Transformer-GAN hybrid. Results are not included, as per your request.

Extended Methodology

To extend our baseline model, we developed and experimented with two additional architectures, each with specialized mechanisms to overcome the limitations observed in standard GAN-based inpainting, e.g., soft textures or lacking spatial consistency. This comparative experiment allowed us to better understand each GAN-based variant's performance limits and trade-offs.

1. DCGAN-Based Encoder-Decoder Generator

We adapted the DCGAN architecture here to support image inpainting. The generator was a typical encoder-decoder setup, processing the masked image and the binary mask. The encoder compressed the visual information into a latent space representation, and the decoder restored the missing content.

For contextual understanding facilitation, we used an extra parallel mask encoder that learned from the structure of the missing patch. Outputs from both encoders were concatenated before decoding to the target image. This helped the model pay more attention to the masked area. Patch-based discriminator was used to test the realism of the output image so that adversarial feedback can be provided to the generator. Adversarial loss as well as pixel-wise L1 reconstruction loss was included during training.

2. BigGAN-Inspired Residual Generator with Attention

Inspired by BigGAN, we added residual learning and attention to make the generator more expressive. The structure had multiple residual blocks on the encoder and decoder sides. Residual blocks helped the model preserve high-frequency information between layers, thereby improving image texture and acuteness.

Half-way through the generator, a self-attention mechanism was added in an attempt to capture global relationships more effectively, particularly for the larger missing regions. The attention block helps the network focus on long-distance relations in the image, e.g., symmetry or texture consistency across the face. This model displayed increased coherence within the restored content, especially for complex structures like eyes or hair. The discriminator was identical in architecture but trained with a cosine learning rate scheduler.

3. Transformer-GAN Hybrid Model

The most recent model fused the pixel-precision of convolutional layers with the global awareness of Transformer blocks. We replaced traditional convolution-only feature extraction with Transformer-based attention modules in the generator. These attention blocks enabled the model to peer beyond local patches and use contextual information across the entire image, leading to more accurate structural restoration.

The generator used a UNet-style encoder-decoder backbone supplemented with residual blocks and Transformer layers. We combined adversarial loss, pixel-wise L1 loss, and perceptual loss (calculated from pretrained VGG16 features) for optimization in loss. We used a PatchGAN-based discriminator that was learned to search for fine-grain realism as well as structure.

By adding perceptual and SSIM-based losses, we tuned the output not only for pixel-level correctness but also for perceptual similarity and visual realism.

The combination of these three models collectively build a progressive improvement pipeline—start from a rudimentary encoder-decoder (DCGAN) towards a dense attention-based generator (BigGAN-inspired), then towards a transformer-GAN blend that excels at retaining both texture and structure. The module-based and comparative approach gave us better insight into how diverse GAN components make contributions towards image inpainting successes on the CelebA set.

Results:

Based on the performance of all three GAN-based inpainting models, we compared their results in terms of three evaluation metrics: L1 Loss, PSNR (Peak Signal-to-Noise Ratio), and SSIM (Structural Similarity Index Measure). These metrics tell us about how precise and visually significant the produced images are compared to the original images.

Model 1: DCGAN-based Encoder-Decoder Model

The initial model used a simple encoder-decoder GAN with mask conditioning. It worked comparatively well for inpainting images, especially when there were small patches of missing data. The final test PSNR was 23.18 dB, and generator loss went down significantly after training to a mean of 3.1971. This model would sometimes not generate fine textures, especially when significant parts of an image were masked.

Model 2: BigGAN-Inspired Residual GAN

The second model improved upon the baseline through the inclusion of residual blocks and attention. This ensured that the generator placed higher importance on related image features. As

a result, it now had much improved visual coherence when generating inpainting results. Performance metrics showed dramatic improvement with L1 Loss being lowered to 0.3581, PSNR rising to 29.57 dB, and SSIM to 0.6590. This means that the second model generated better and more structurally accurate images than the baseline model.

FINAL Model: Transformer-GAN Hybrid

Our final model combined the benefit of both GANs and Transformers. The GAN helped in creating high-level texture, and the Transformer could learn to incorporate the global context of the image. The model yielded the best performance among the three. The average L1 Loss reduced significantly to 0.0265, PSNR rose to 26.19 dB, and SSIM had a large value of 0.8794, which is suitable for high similarity to the original images.

We also graphed the distribution of these metrics. The histogram plots revealed that most images had low L1 Loss and high SSIM, while PSNR was in the range of 25–27 dB. Even the box plots confirmed for us that our model's predictions were consistent and reliable. Also, the training loss curve showed smooth convergence of both generator and discriminator losses. Lastly, the scatter plot of PSNR vs. SSIM indicated an extremely high positive correlation that proved visual quality was also enhanced when structural similarity increased.

Plots and visualizations:




Fig:1




Fig : 2




Fig : 3





Fig : 4



Trello Work distribution screen shots :

Key	Summary	Status	Assignee
GI-2	Define project objectives and research problem	DONE	TM Taran Mamidala
GI-3	Study existing image inpainting techniques (DeepFill v2,...	DONE	SS SAI LAKSHMI KALYAN VEERA SAIVA
GI-4	Research Transformer-based inpainting models	DONE	SJ Sai Abhisekh Reddy Jupalli
GI-5	Finalize project methodology and architecture	DONE	TM Taran Mamidala
GI-6	Identify required datasets and preprocess them	DONE	SJ Sai Abhisekh Reddy Jupalli
GI-7	Set up project repository (GitHub) and documentation	DONE	SS SAI LAKSHMI KALYAN VEERA SAIVA
GI-8	Develop Transformer-based encoder	DONE	SJ Sai Abhisekh Reddy Jupalli
GI-9	Implement GAN-based decoder	DONE	TM Taran Mamidala
GI-10	Integrate PatchGAN discriminator	DONE	SS SAI LAKSHMI KALYAN VEERA SAIVA
GI-11	Implement loss functions (adversarial loss, SSIM, perce...	DONE	SJ Sai Abhisekh Reddy Jupalli
GI-12	Set up the self-supervised pretraining pipeline	DONE	TM Taran Mamidala
GI-13	Fine-tune model using inpainting datasets	DONE	SJ Sai Abhisekh Reddy Jupalli
GI-14	Implement evaluation metrics (PSNR, SSIM, FID, LPIPS)	DONE	SS SAI LAKSHMI KALYAN VEERA SAIVA
GI-15	Train and optimize the model	DONE	TM Taran Mamidala
GI-16	Test model performance on different datasets	DONE	TM Taran Mamidala
GI-17	Compare results with DeepFill v2, LaMa, and CoMod-G...	DONE	SJ Sai Abhisekh Reddy Jupalli
GI-18	Fine-tune hyperparameters for better results	DONE	TM Taran Mamidala
GI-19	Generate sample inpainted images and analyze outputs	DONE	TM Taran Mamidala
GI-20	Prepare the final report and documentation	DONE	SJ Sai Abhisekh Reddy Jupalli
GI-21	Create visuals, graphs, and result analysis	DONE	SJ Sai Abhisekh Reddy Jupalli
GI-22	Submit project and prepare for presentation	DONE	SS SAI LAKSHMI KALYAN VEERA SAIVA
GI-23	Review project feedback and final refinements	DONE	TM Taran Mamidala

Chart



Deployed Application Screen shots:



Controls

Upload Image

Drag and drop file here

Limit 200MB per file • PNG, JPG, JPEG

Browse files

000019.jpg

7.7KB

×

Mask Mode

Random

Freehand

Random Mask Size

64

32

128

Brush Width

20

5

50

Run inpainting

Inpainted Output



Controls

Upload Image

Drag and drop file here

Limit 200MB per file • PNG, JPG, JPEG

Browse files

000019.jpg

7.7KB

×

Mask Mode

Random

Freehand

Random Mask Size

64

32

128

Brush Width

20

5

50

Run inpainting

000019.jpg

7.7KB


×


0.0102

0.9563

26.10 dB

Error Distribution





Controls

Upload Image

Drag and drop file here
Limit 200MB per file • PNG, JPG, JPEG

Browse files

000012.jpg
5.0KB

Mask Mode

☐ Random

☒ Freehand


Random Mask Size

32 64 128

Brush Width

5 20 50


Run inpainting



Draw Mask



Masked Input



Controls

Upload Image

Drag and drop file here
Limit 200MB per file • PNG, JPG, JPEG

Browse files

000012.jpg
5.0KB

Mask Mode

☐ Random

☒ Freehand

Random Mask Size

32 64 128

Brush Width

5 20 50

Run inpainting

Masked Input



Controls

Upload Image

Drag and drop file here
Limit 200MB per file • PNG, JPG, JPEG

Browse files

000012.jpg
5.0KB

Mask Mode

☐ Random

☒ Freehand

Random Mask Size

32 64 128

Brush Width

5 20 50

Run inpainting

Inpainted Output



References:

1. **Pathak et al. (2016).** “Context Encoders: Feature Learning by Inpainting.” *Proceedings of CVPR 2016*.
2. **Iizuka et al. (2017).** “Globally and Locally Consistent Image Completion.” *ACM Transactions on Graphics (SIGGRAPH Asia 2017)*.
3. **Liu et al. (2018).** “Image Inpainting for Irregular Holes Using Partial Convolutions.” *European Conference on Computer Vision (ECCV 2018)*.
4. **Yu et al. (2019).** “Free-Form Image Inpainting with Gated Convolution.” *Proceedings of ICCV 2019*.
5. **Nazeri et al. (2019).** “EdgeConnect: Generative Image Inpainting with Adversarial Edge Learning.” *Proceedings of ICCVW 2019*.
6. **Liu et al. (2015).** “Deep Learning Face Attributes in the Wild.” *Proceedings of ICCV 2015*. – Introduces the CelebA dataset used for training and evaluation.
7. **Paszke et al. (2019).** “PyTorch: An Imperative Style, High-Performance Deep Learning Library.” *NeurIPS 2019 Workshop*.
8. **Streamlit Inc. (2019).** “Streamlit: Turn Python Scripts into Shareable Web Apps in Minutes.”
9. **Wang et al. (2004).** “Image Quality Assessment: From Error Visibility to Structural Similarity.” *IEEE Transactions on Image Processing*.
10. **Hore & Ziou (2010).** “Image Quality Metrics: PSNR vs. SSIM.” *Proceedings of ICPR 2010*.